




ADVANCED AVF RSS & FDIR

Yahui Cao

Agenda



AVF Introduction

RSS Introduction

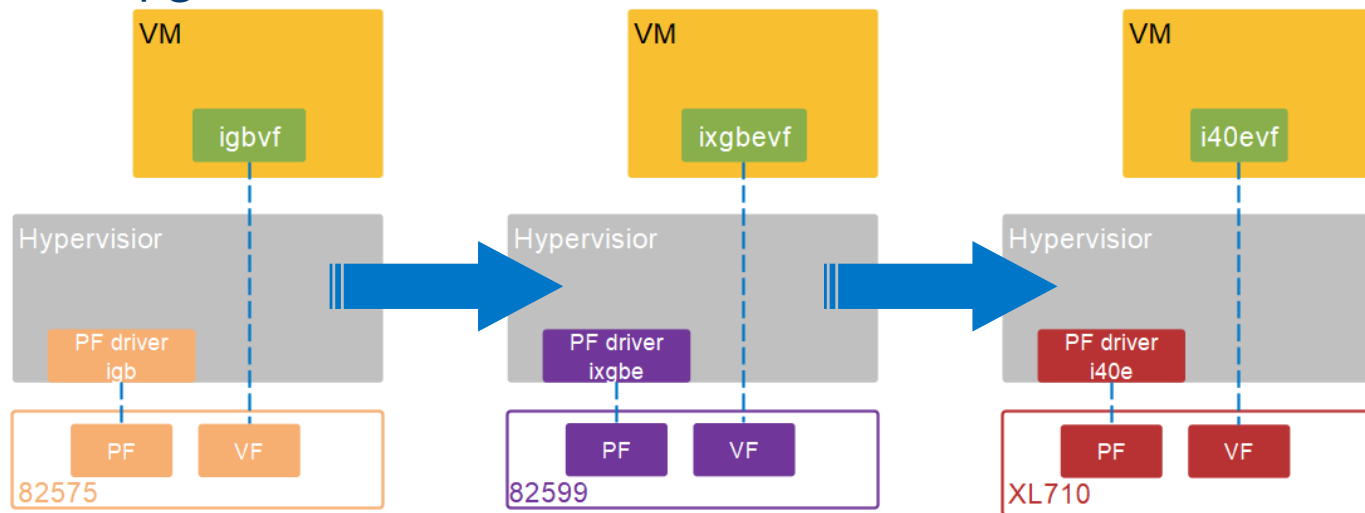
FDIR Introduction

Virtual Channel Introduction

AVF Intro

- Usually different SRIOV Virtual Function has different VF driver


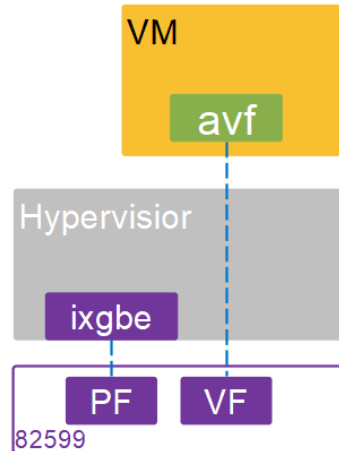
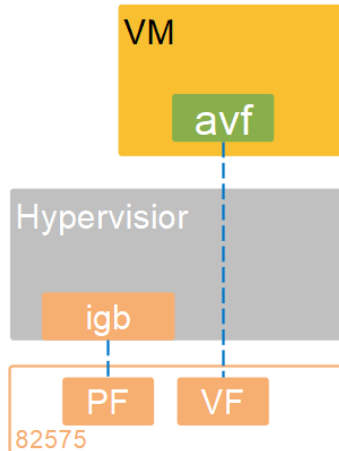
HW Upgrade



AVF Intro

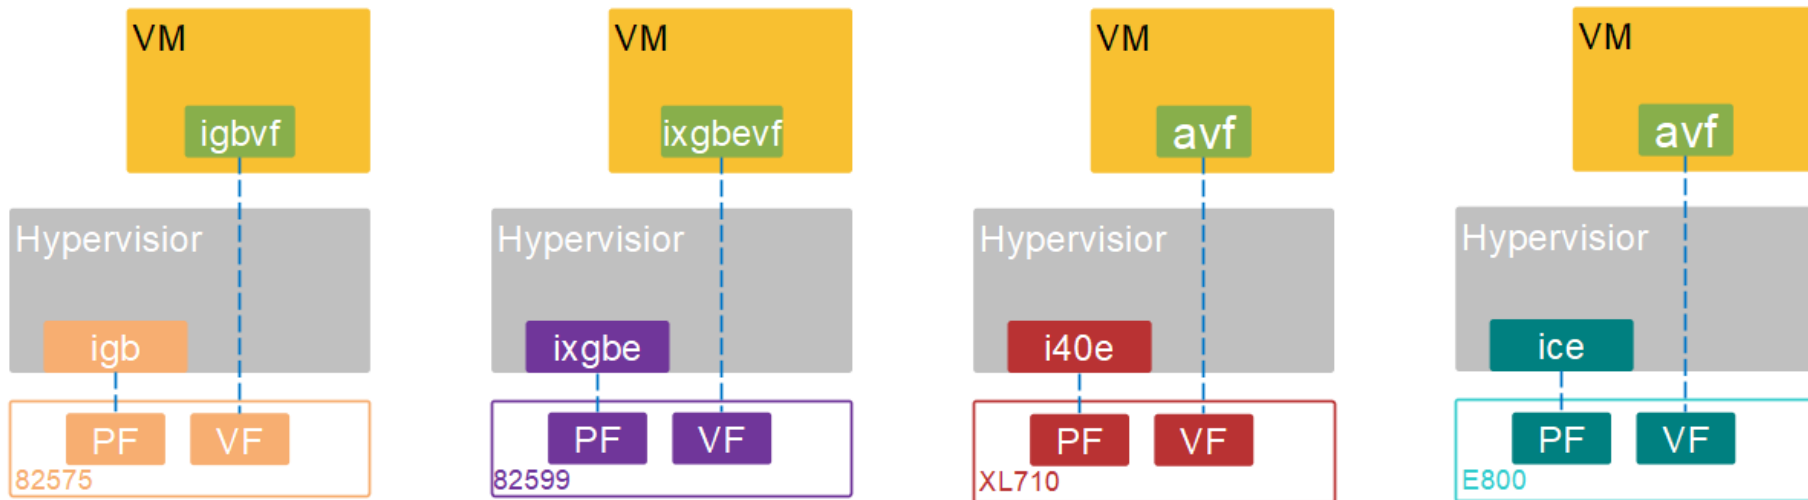
- AVF(Adaptive Virtual Function) is a specification:
 - Let one VF driver to drive different intel NIC
 - Let VM image run on the new NIC without changes.

Ideal AVF




AVF Intro

Real World AVF




AVF Intro



兼容性。

- How to support NIC's new feature: Base + Advanced
- Base mode support:
 - Device ID, mailbox
 - Checksum, TSO offload
 - Tx/RX, Multiqueue, MSIX
- Advanced feature:
 - Depend on NIC's capability
 - Need negotiation.

AVF Intro



The following minimal features are supported by the AVF driver:

1. A mailbox to the PF driver
2. 4 Rx and Tx queue pairs
3. 5 MSI-X interrupts
4. Per vector interrupt moderation
5. An RSS table of 64 entries that can point to up to 4 queues
6. Basic Rx and Tx offloads (checksum and TSO for non-tunneled packets)

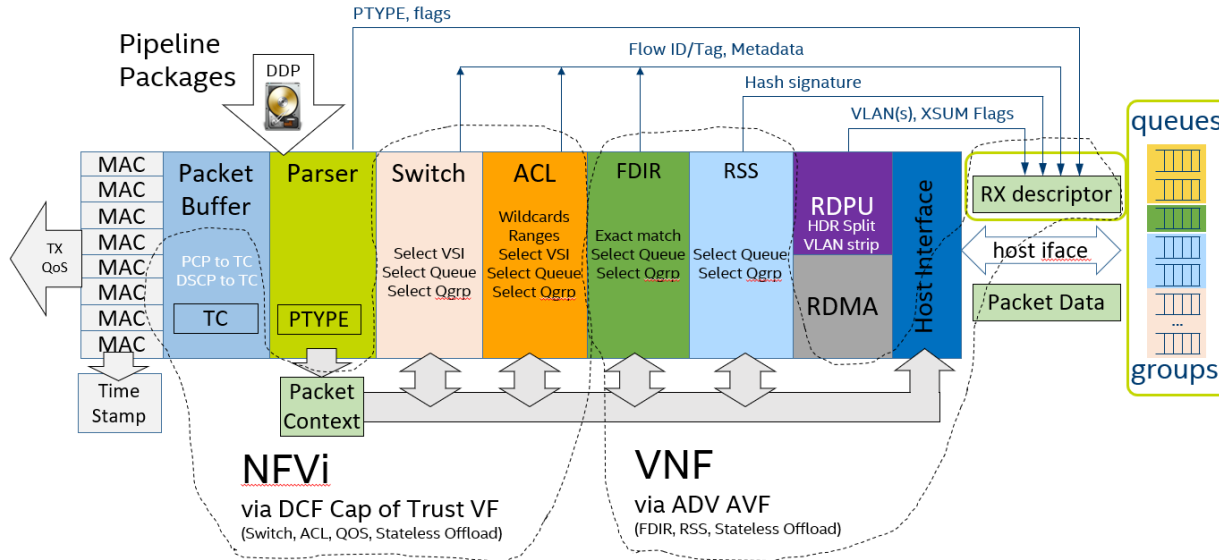
The following features are optional and might be exposed by the hardware:

1. More queues or more interrupts.
2. Extension of the RSS table so that it can point to a larger number of queues
3. RDMA support
4. VF VLAN trunking
5. VF promiscuous
6. Negotiate header-split
7. Negotiate Tx checksum/TSO for tunneled packets

ADV Intro

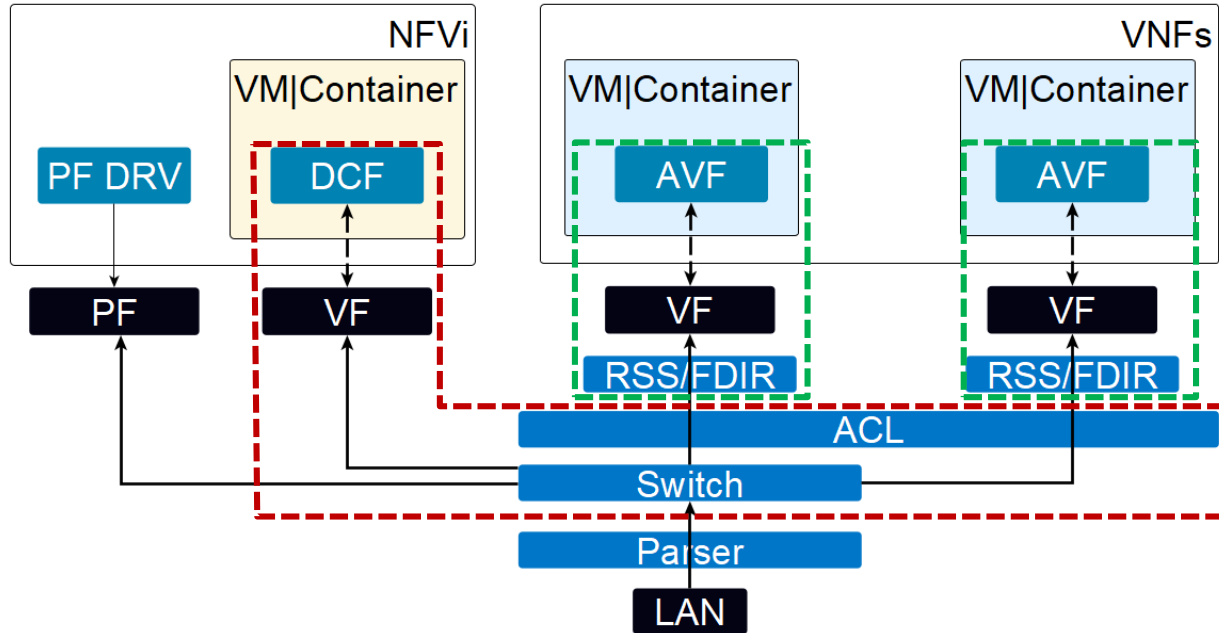
➤ CVL RSS & FDIR capability mapping into AVF

Intel® Ethernet 800 Series Programmable Pipeline




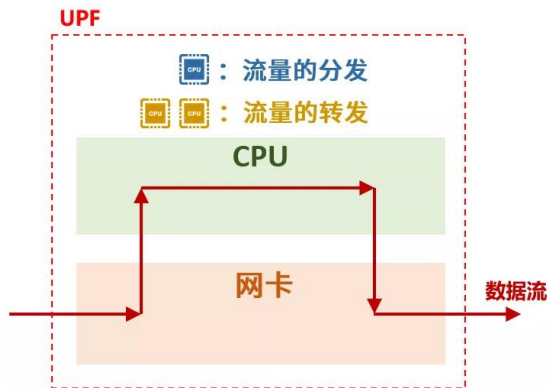
ADV Intro

➤ CVL RSS & FDIR capability mapping into AVF




AVF Use Case

- 5G UPF use 1/3 cores for traffic distribution. 需要1/3 core进行流量分发。
- Use DCF + AVF can offload cpu



Agenda



AVF Introduction


RSS Introduction

FDIR Introduction

Virtual Channel Introduction

RSS Intro


In early days, all network packets are forwarded to the single core.



早期，一个queue，一个core。

RSS Intro


Here, all packets are forwarded only to the specific one process core. Network does not benefit from multicore.



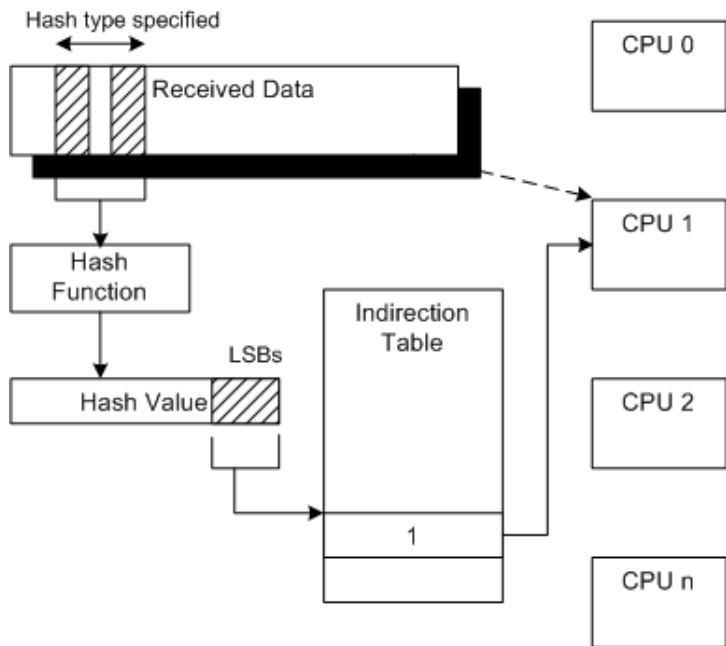
package是到queue，不是到core吧

RSS Intro

Then **Microsoft** introduces RSS to do load balancing.



RSS Intro



- Early days, NIC has only 1 queue. RSS is done by cpu cores on Windows.

只有一个queue怎么做RSS?

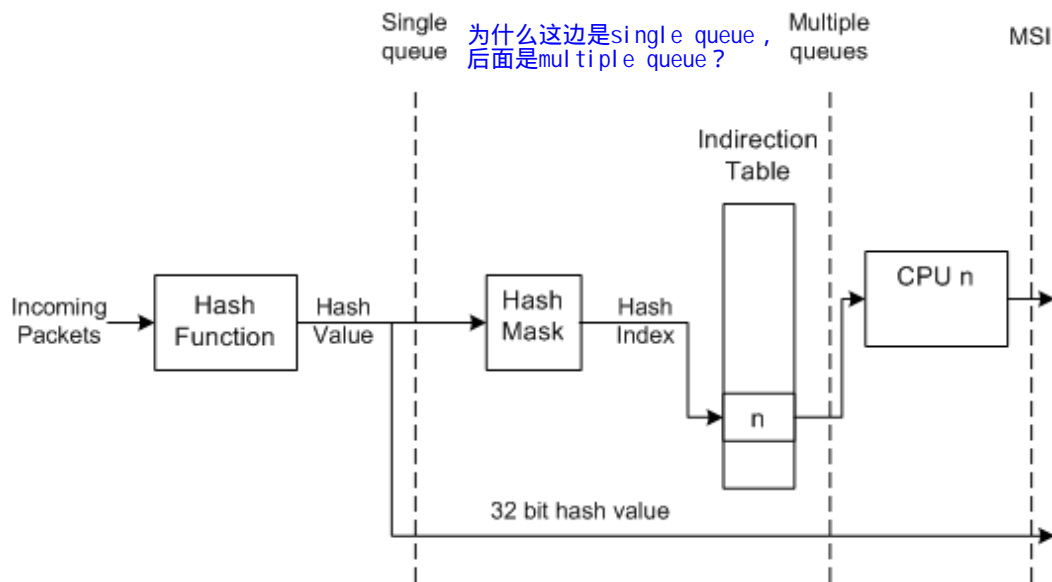
- Linux RSS software implementation:
Receive Packet Steering
Receive Flow Steering

RSS Intro

➤ NIC support for RSS:

1. Hash for Single queue
2. Hash for multi queue
3. MSI/MSIX

算hash, 查表, 分发。




RSS Use Case

- Default setup:
5-tuple

- Business specific distribution:
PPPoE, GTP-U, etc

```
flow create 0 ingress pattern eth / ipv4 / udp / gtpu / gtp_psc / ipv4 / end  
actions rss types ipv4 l3-src-only end key_len 0 queues end / end
```

Agenda



AVF Introduction


RSS Introduction

FDIR Introduction

Virtual Channel Introduction

FDIR Intro


However we can't let data go to specific core, which still destroys multicore performance



FDIR Intro


精准导流。细粒度优化。

Then Intel Flow Director is coming! FDIR can send incoming packets directly to the right core.



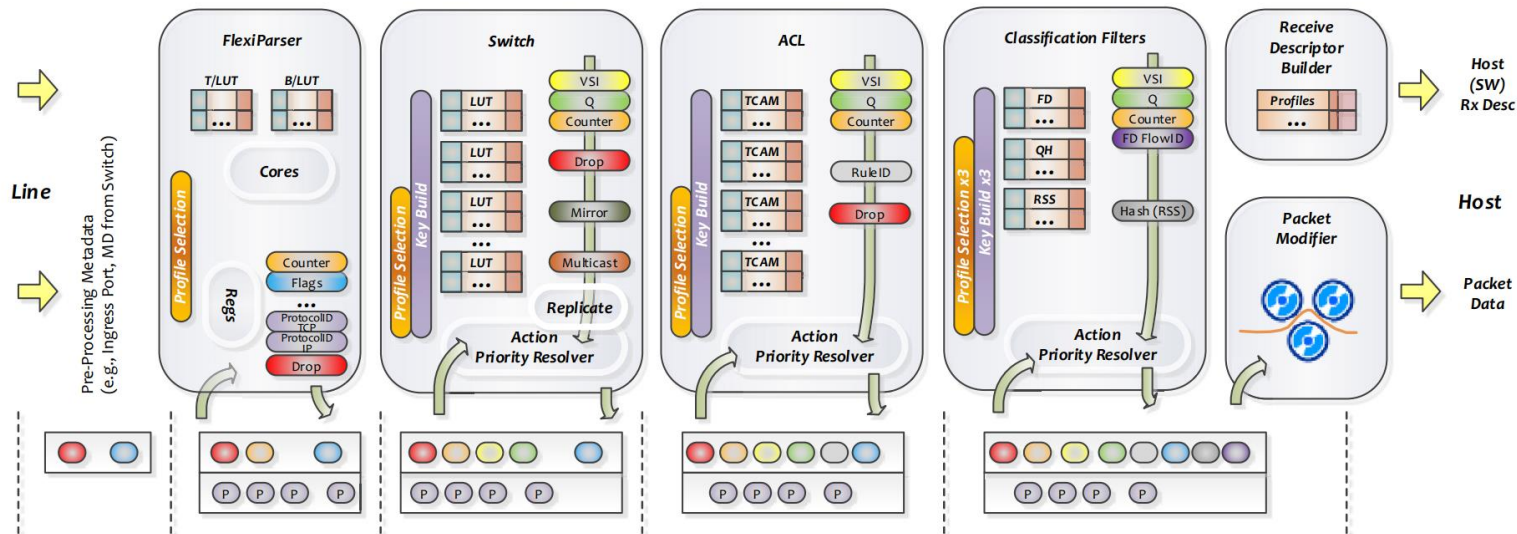
FDIR Intro

Flow Director has a match table to send matched packets to the desired process core.



FDIR Intro

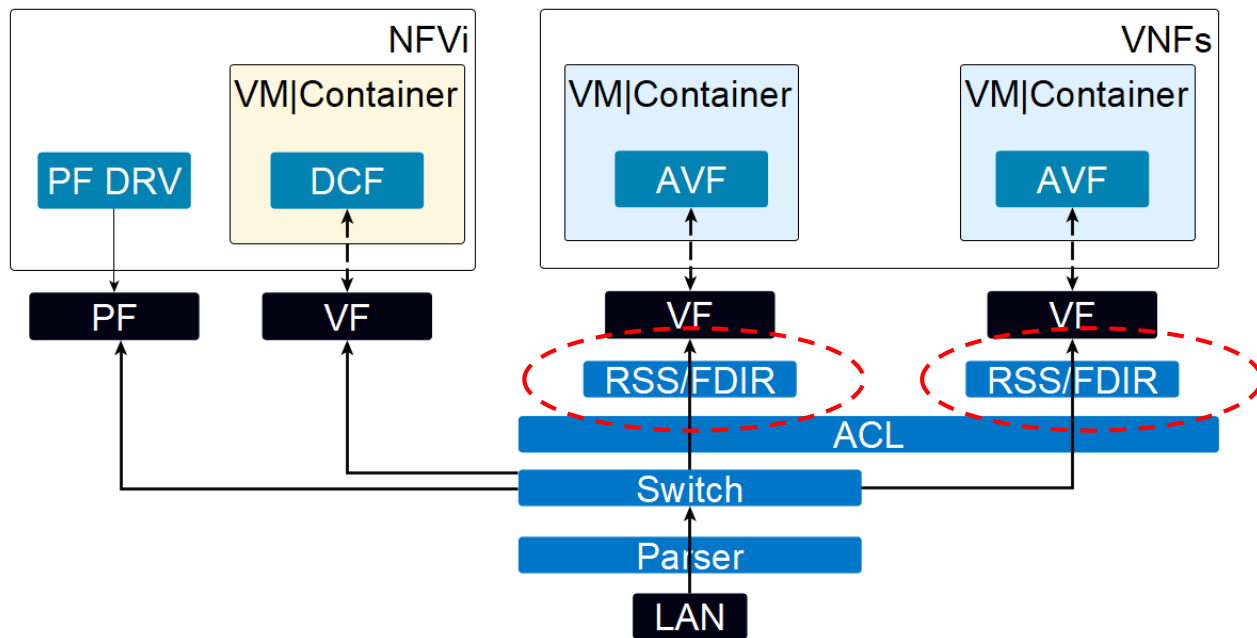
Where is FD in CVL packet processing pipeline?



不同颜色的圆圈代表不同的action (e.g. switch: to Q0, ACL: to Q1) , 它们有不同的priority.

ADV Intro

Where is FD in CVL packet processing pipeline?

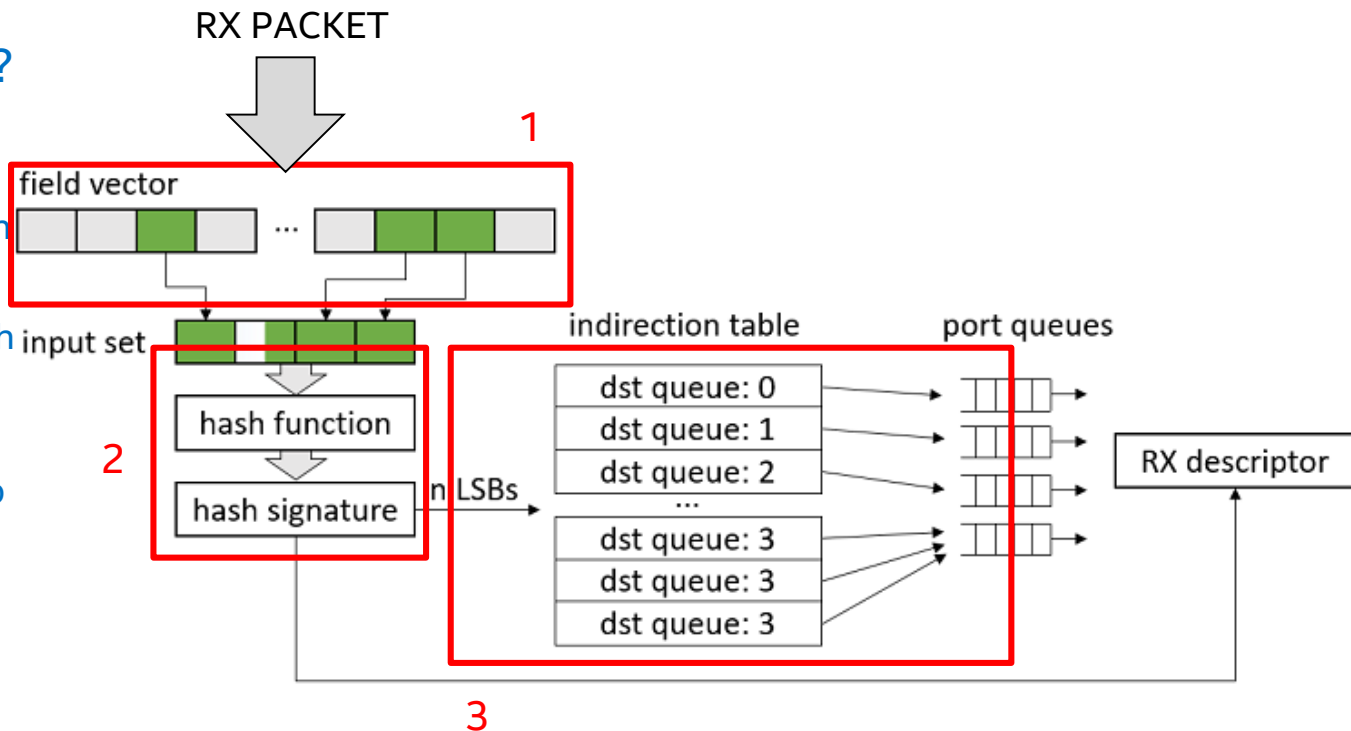


先FDIR, 再RSS。
如果FDIR命中, 则不会再做
后面的匹配。

FDIR Intro

How does FDIR work?

1. Field vector extracts rx packets' specific domain as input set
2. Input set generates hash signature by hash function
3. Hash signature looks up matched item in the indirection table and then do forwarding

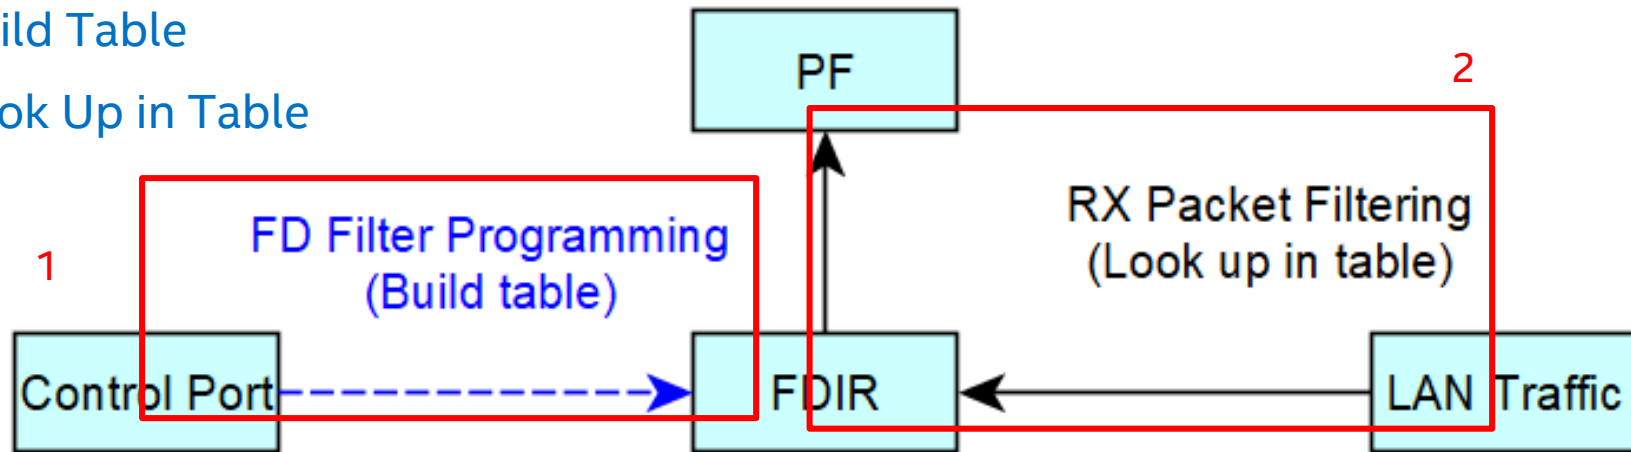


FDIR Intro

How FDIR function block is designed ?

1. Build Table

2. Look Up in Table

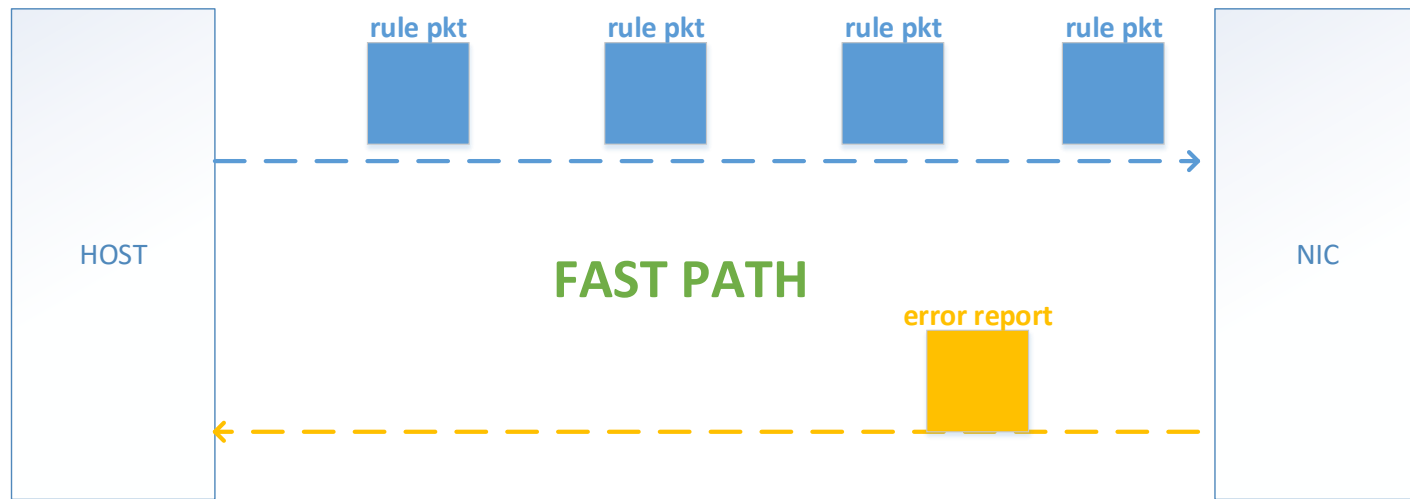


FDIR Intro

通过package写表。
写表的过程其实就是收发包的过程。

How to build FDIR look up table?

By FDIR TX/RX queues !



Picture From DPDK Summit North America 2019, *Rte_flow Optimization in i40e Driver* - Chenmin Sun, Intel

FDIR Intro

1. Enable FDIR queues:

- a) Request for VSI
- b) Request for TX/RX queues
- c) Request for TX/RX irq
- d) Prepare Ring buffer

2. Allocate FDIR rule space

3. Build FDIR rule table

how?

Resource	800 Series	700 Series
Rx Queues	2048	1536
Tx Queues	16K	1536
Tx Completion Queues	512	N/A
Tx Doorbell Queues	256	N/A
MSIX Vectors	2048	1168
RSS	768 (64), 16 (512), 8 (2K)	128 (64)
FD Filter	16K	8K
VSI	768	384
SR-IOV VF	256	128
Ternary Classification	40bits x 8K TCAM	N/A
Binary Classification	80bits x 32K CAM	1K MAC, 512 VLAN
VEB	256 (32 w/ Stats.)	16

8000条
rule->16000条
rule。

FVL vs CVL resource table

CVL Flow Director Enablement

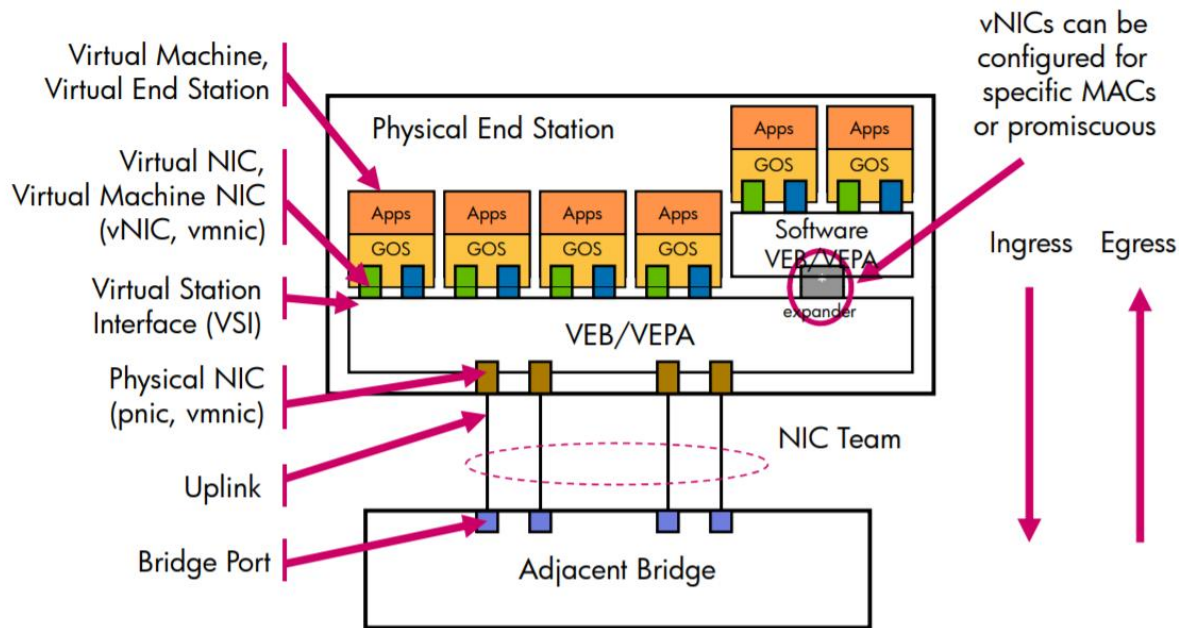
To enable FDIR queues

a) Request for VSI

1. What is VSI?
Virtual Station Interface

An internal point-to-point Ethernet connection which connects an Ethernet port of a VEB or VEPA to a vNIC

2. VSI type in CVL:
VF, VMDq2, PF, EMP/MNG

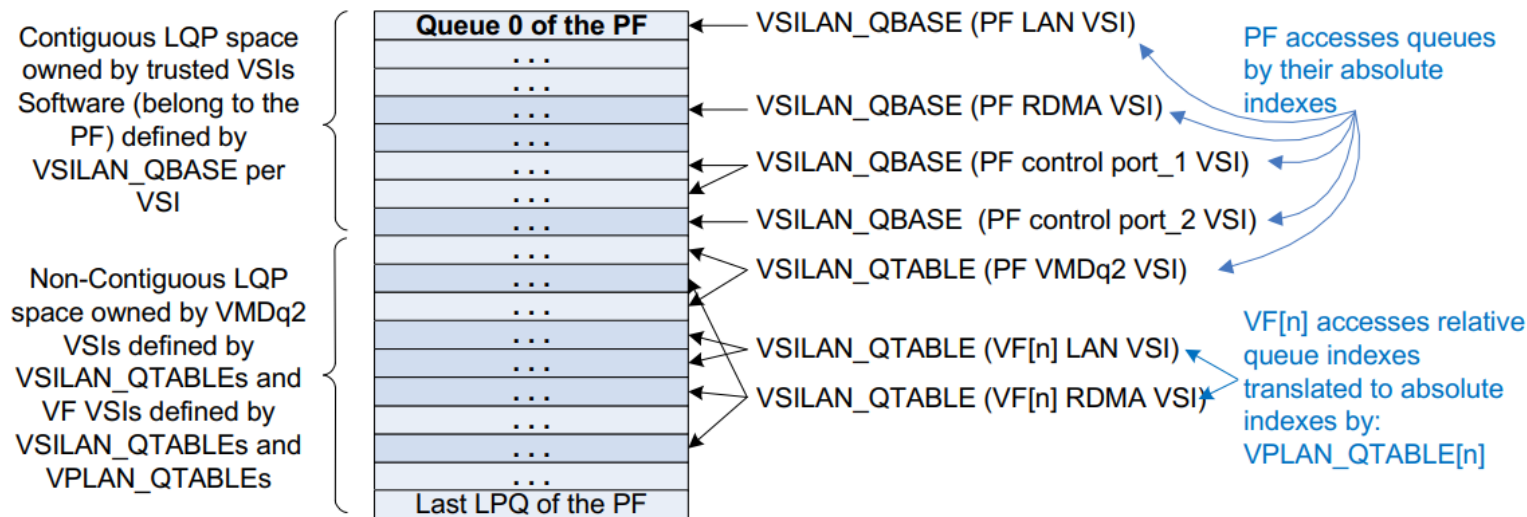


CVL Flow Director Enablement

To enable FDIR queues:

b) Request for TX/RX queue


Queues can be requested in Contiguous/Scattered mode



CVL Flow Director Enablement

To enable FDIR queues:

d) Prepare Ring Buffer



CVL Flow Director Enablement

Allocate FDIR rule space:

How is FD rule space managed?

1. 16K filter rules are shared among PFs by:

Unique (“guaranteed”) 独占

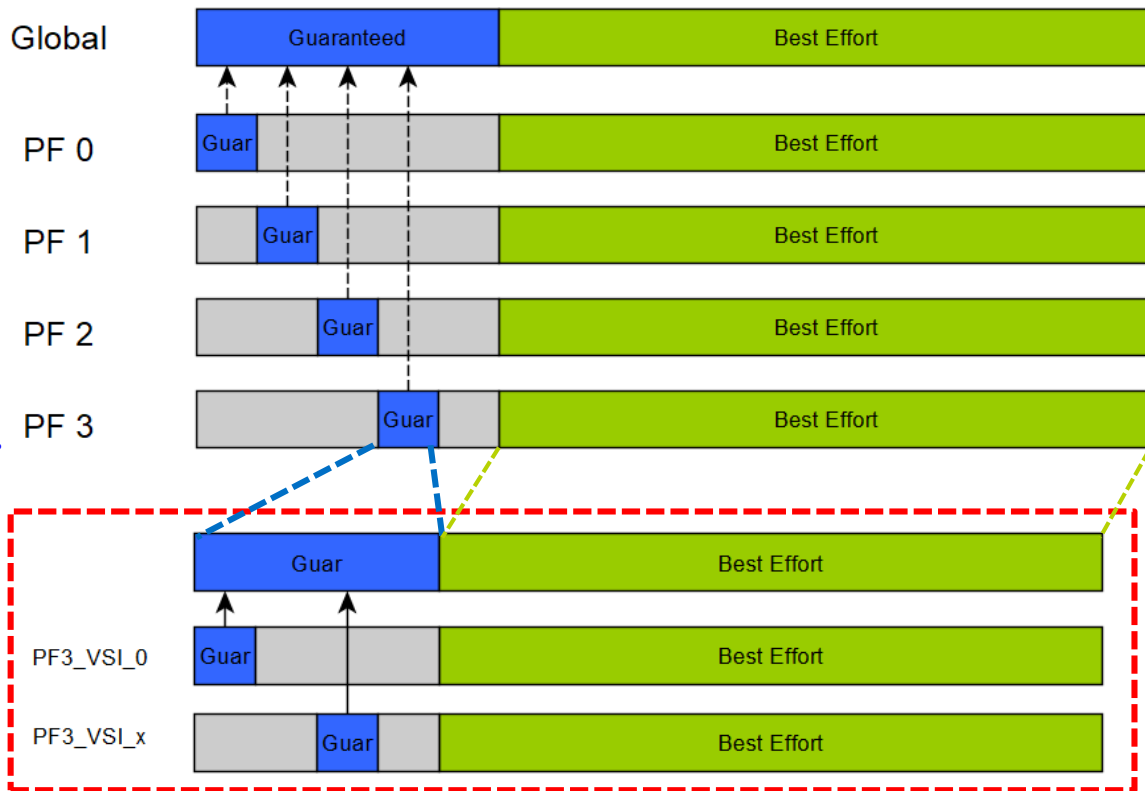
+

Shared (“best effort”) 共享, 先到先得。

e.g. In 25Gx4 Card:
guaranteed = 512,
best_effort =

$$16K - 512 * \text{Active_PF}$$

目前, vf只用best effort里的资源。



CVL Flow Director Enablement

How to Build FDIR rule table?

1. Initialize Field Vector Word

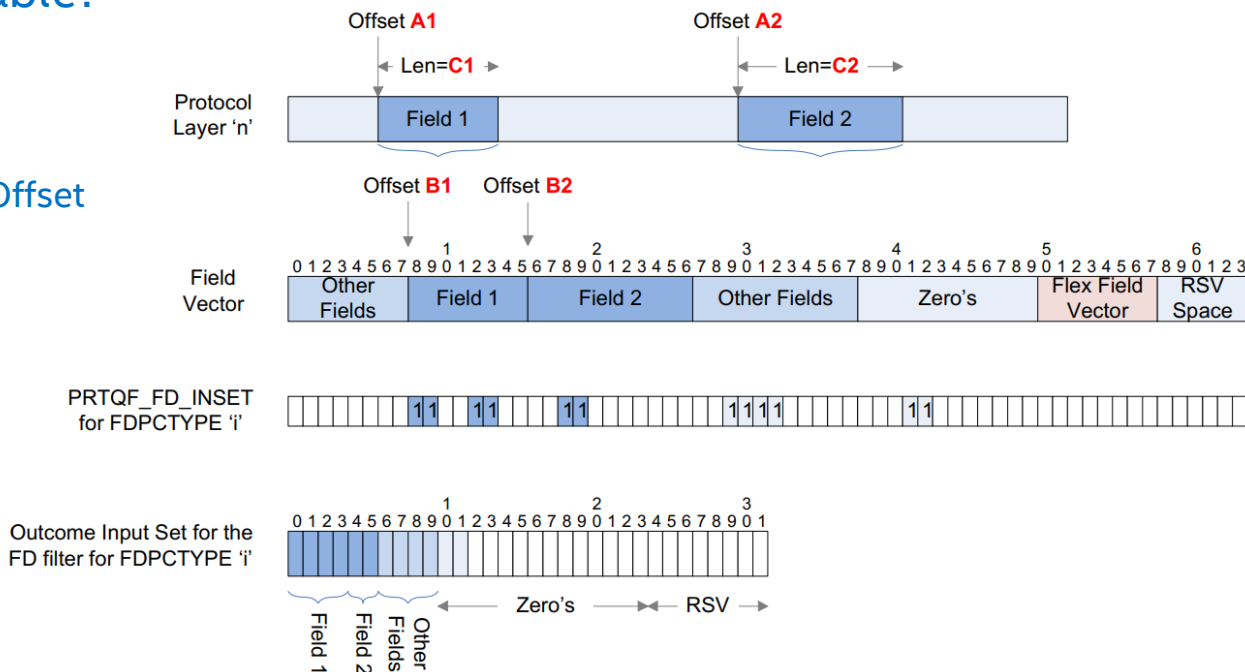
What is Field Vector ?

A combination set of Word Offset
and Word Length

CVL support up to 24 words
matching

The field vector specifies which packet
header fields are considered when creating
FDIR filters,

字段匹配：需要知道字段在以太网包里的偏移和长度。



CVL Flow Director Enablement

匹配destination mac需要3个word

How to Build FDIR rule table?

1. Initialize Field Vector Word


OFFSET= 0, 2, 4



What is Field Vector ?
A combination set of Word Offset
and Word Length

OFFSET = X

OFFSET = X+8



CVL support up to 24 words
matching

OFFSET = Y, Y+2

OFFSET = Y+4, Y+6



```
ice_fdir_filter.c
ice_fdir_parse_pattern(){} //这个过程可以看看
const struct rte_flow_item_ipv4 *ipv4_spec, *ipv4_mask;
```

CVL Flow Director Enablement

How to Build FDIR rule table?

1. Initialize Field Vector Word

FV is initialized by creating Profile

E.g.

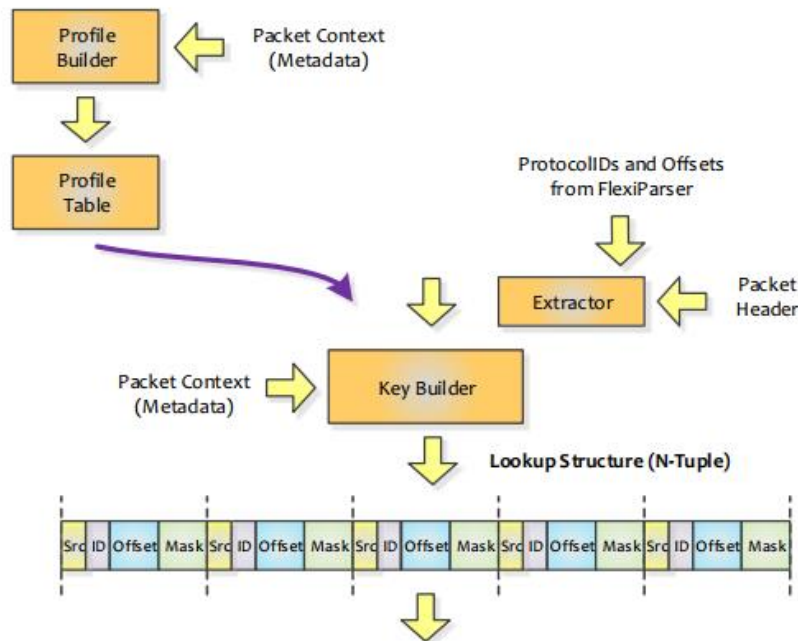
IPV4+TCP/UDP/SCTP

IPV6+TCP/UDP/SCTP

IPV4+UDP+VXLAN+MAC+IPV4+TCP
/UDP/SCTP

IPV4+UDP+GTPU+PDU+IPV4+TCP/
UDP/SCTP

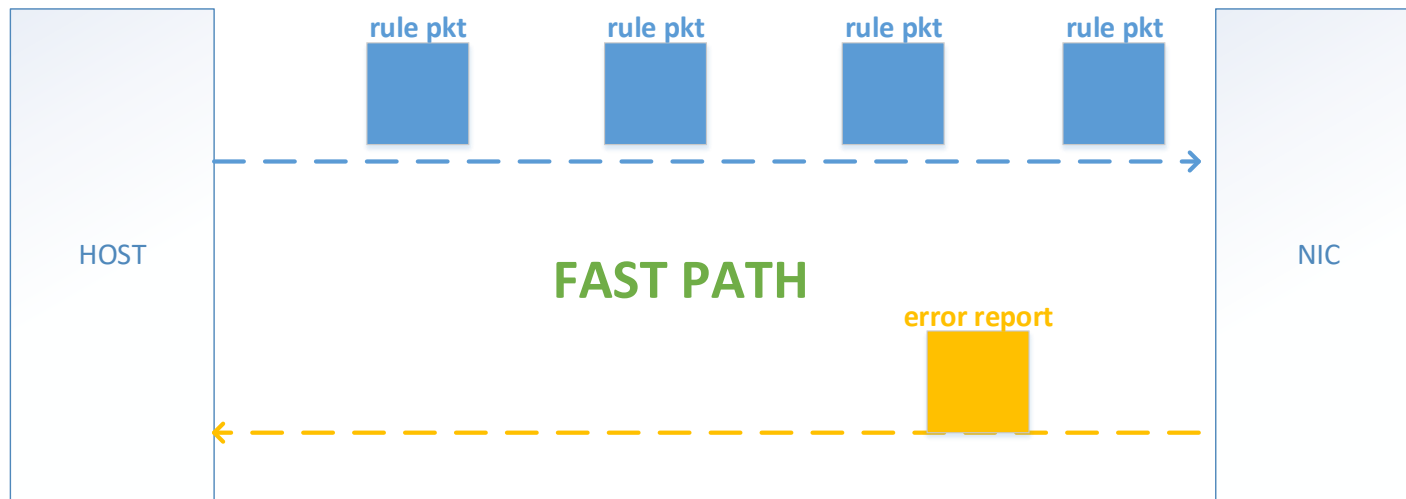
创建profile的过程就是在初始化field vector。



CVL Flow Director Enablement

How to Build FDIR rule table?

1. Create Profile
2. Program the input set value and table by FDIR TX queue.

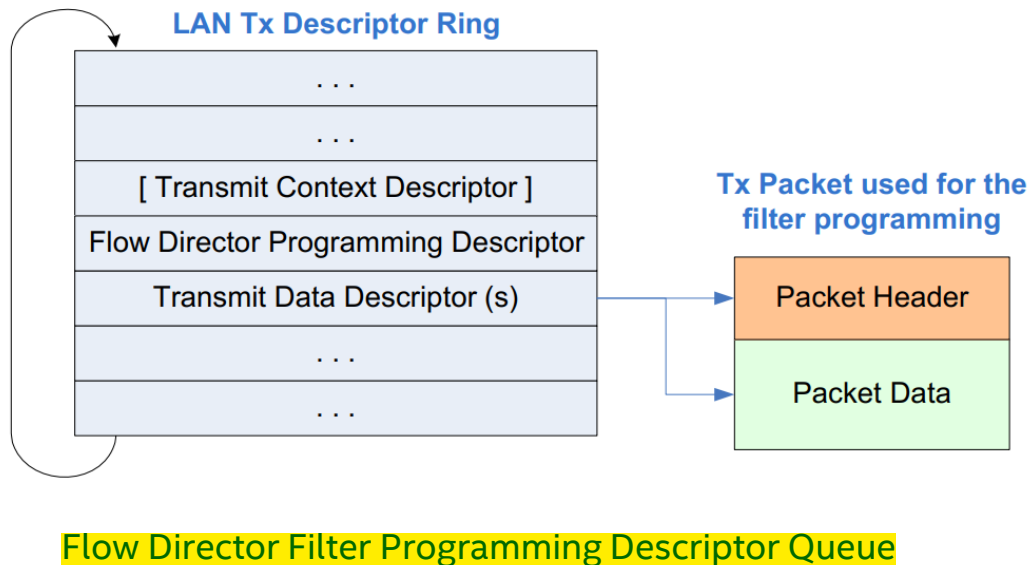


Picture From DPDK Summit North America 2019, *Rte_flow Optimization in i40e Driver* - Chenmin Sun, Intel

CVL Flow Director Enablement

How to Build FDIR rule table?

1. Create Profile
2. Program the input set value and table by FDIR TX queue.
 - a) Get free queue descriptor



CVL Flow Director Enablement

填包。

How to Build FDIR rule table?

1. Create Profile
2. Program the input set value and table by FDIR TX queue.
 - a) Get free queue descriptor
 - b) Fill LAN+FDIR descriptor

[illegible]

FD Filter Programming Descriptor

[illegible]

Transmit Data Descriptor

CVL Flow Director Enablement

How to Build FDIR rule table?

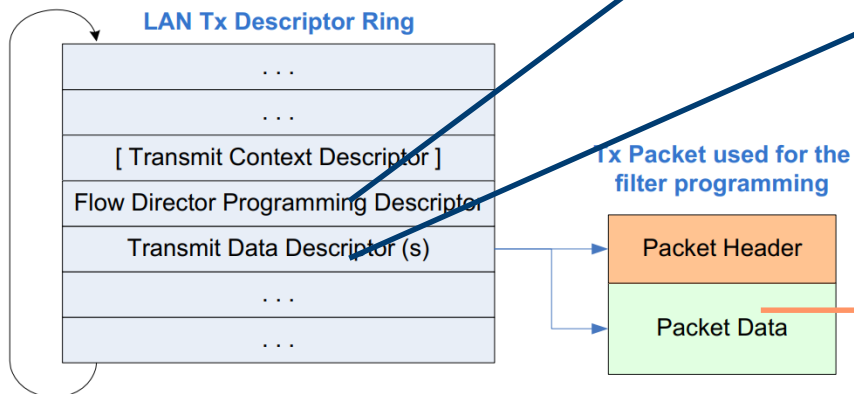
1. Create Profile
2. Program the input set value and table by FDIR TX queue.
 - a) Get free queue descriptor
 - b) Fill LAN+FDIR descriptor
 - c) Prepare FDIR dummy packet

```
static const u8 ice_fdir_tcpv4_pkt[] = {  
    0x00, 0x00, 0x00, 0x00, 0x00, 0x00, 0x00, 0x00,  
    0x00, 0x00, 0x00, 0x00, 0x08, 0x00, 0x45, 0x00,  
    0x00, 0x28, 0x00, 0x01, 0x00, 0x00, 0x40, 0x06,  
    0x00, 0x00, 0x00, 0x00, 0x00, 0x00, 0x00, 0x00,  
    0x00, 0x00, 0x00, 0x00, 0x00, 0x00, 0x00, 0x00,  
    0x00, 0x00, 0x00, 0x00, 0x00, 0x00, 0x50, 0x00,  
    0x20, 0x00, 0x00, 0x00, 0x00, 0x00  
};
```

CVL Flow Director Enablement

How to Build FDIR rule table?

1. Create Profile
2. Program the input set value and table by FDIR TX queue.

[illegible][illegible]

```
static const u8 ice_fdir_tcpv4_pkt[] = {
    0x00, 0x00, 0x00, 0x00, 0x00, 0x00, 0x00, 0x00,
    0x00, 0x00, 0x00, 0x00, 0x08, 0x00, 0x45, 0x00,
    0x00, 0x28, 0x00, 0x01, 0x00, 0x00, 0x40, 0x06,
    0x00, 0x00, 0x00, 0x00, 0x00, 0x00, 0x00, 0x00,
    0x00, 0x00, 0x00, 0x00, 0x00, 0x00, 0x00, 0x00,
    0x00, 0x00, 0x00, 0x00, 0x00, 0x00, 0x50, 0x00,
    0x20, 0x00, 0x00, 0x00, 0x00, 0x00
};
```

CVL Flow Director Enablement

How to Build FDIR rule table?

1. Create Profile
2. Program the input set value and table by FDIR TX queue.
 - a) Get free queue descriptor
 - b) Fill LAN+FDIR descriptor
 - c) Prepare FDIR dummy packet
 - d) Notify HW (Doorbell) and HW Write Back
3. Check programming status (Optional)
4. Check rule table space and size (Highly recommend)

CVL Flow Director Debug

How to debug FDIR rule ?

(E.g. pattern eth / ipv4 dst is 192.168.0.2 / end)

1. Check if Profile is created successfully and correctly for both control port and PF port.

Notes:

What's the difference between control port and main port in FDIR ?

Control port in FDIR is responsible for building rule table by programming dummy packet.
PF Port in FDIR is responsible for filtering packets by looking up rule table.

FDIR set rule steps

1. Create Profile

2. Set rule through Flow Director TX/RX queue
 - a) Get free queue descriptor
 - b) Fill LAN+FDIR descriptor
 - c) Prepare FDIR dummy packet
 - d) Notify HW (Doorbell) and HW Write Back
3. Check programming status (Optional)
4. Check rule table space and size (Highly recommend)

CVL Flow Director Debug

How to debug FDIR rule ?

(E.g. pattern eth / ipv4 dst is 192.168.0.2 / end)

1. Check if Profile is created successfully and correctly: Bind correct ptype ?

```
ptype 22: MAC_IPV4_FRAG
ptype 23: MAC_IPV4_PAY
ptype 24: MAC_IPV4_UDP_PAY
ptype 25: PT_RESERVED25
ptype 26: MAC_IPV4_TCP
ptype 27: MAC_IPV4_SCTP
ptype 28: MAC_IPV4_ICMP
ptype 29: MAC_IPV4_IPV4_FRAG
```

PTG_IPV4_OTHER MAC_IPV4_FRAG MAC_IPV4_PAY MAC_IPV4_ICMP MAC_IPV4_TUN_PAY

443	24577	PTG_TUN_GTPU_INNER_IPV6_TCP	Ignore
444	24577	PTG_TUN_GTPU_INNER_IPV6_UDP	Ignore
445	24577	PTG_TUN_GTPU_INNER_IPV6_OTHER	Ignore
446	24577	PTG_TUN_GTPU_INNER_IPV4_TCP	Ignore
447	24577	PTG_TUN_GTPU_INNER_IPV4_UDP	Ignore
448	24577	PTG_TUN_GTPU_INNER_IPV4_OTHER	Ignore
449	24577	PTG_PPP0E_IPV4_TCP	Ignore
456	24577	PTG_PPP0E_IPV4_UDP	Ignore
505	24577	PTG_PPP0E_IPV4_OTHER	Ignore
506	24577	PTG_DEFAULT	Ignore
507	24577	PTG_IPV4_SCTP	Ignore
508	24577	PTG_IPV4_TCP	Ignore
509	24577	PTG_IPV4_UDP	Ignore
510	24577	PTG_IPV4_OTHER	Ignore

CVL Flow Director Debug

How to debug FDIR rule ?

(E.g. pattern eth / ipv4 dst is 192.168.0.2 / end)

1. Check if Profile is created successfully and correctly: Bind correct field vector ?

Thanks for Zhang Qi's great package analysis utility !

17	UDP_DEPTH0	12
00	PD_META_255	511
01	PD_META_255	511
02	PD_META_255	511
03	PD_META_255	511
04	PD_META_255	511
05	PD_META_255	511
06	PD_META_255	511
07	PD_META_255	511
08	PD_META_255	511
09	PD_META_255	511
0A	PD_META_255	511
0B	PD_META_255	511
0C	PD_META_255	511
0D	PD_META_255	511
0E	PD_META_255	511
0F	PD_META_255	511
10	PD_META_255	511
11	PD_META_255	511
12	PD_META_255	511
13	PD_META_255	511
14	IPV4_DEPTH0	14
15	IPV4_DEPTH0	12
16	IPV4_DEPTH0	18
17	IPV4_DEPTH0	16
00	PD_META_255	511

CVL Flow Director Debug

How to debug FDIR rule ?

(E.g. pattern eth / ipv4 dst is 192.168.0.2 / end)

1. Check if Profile is created successfully and correctly
2. Check if dummy packet contains the correct training data

FDIR set rule steps

1. Create Profile
2. **Set rule through Flow Director TX/RX queue**
 - a) **Get free queue descriptor**
 - b) **Fill LAN+FDIR descriptor**
 - c) **Prepare FDIR dummy packet**

```
(gdb) x/64xb pkt
0x17fbc0700: 0x00 0x00 0x00 0x00 0x00 0x00 0x00 0x00 0x00
0x17fbc0708: 0x00 0x00 0x00 0x00 0x00 0x08 0x00 0x45 0x00
0x17fbc0710: 0x00 0x14 0x00 0x00 0x40 0x00 0x00 0x00 0x00
0x17fbc0718: 0x00 0x00 0xc0 0xa8 0x00 0x02 0x00 0x00 0x00
0x17fbc0720: 0x00 0x00 0x00 0x00 0x00 0x00 0x00 0x00 0x00
0x17fbc0728: 0x00 0x00 0x00 0x00 0x00 0x00 0x00 0x00 0x00
0x17fbc0730: 0x00 0x00 0x00 0x00 0x00 0x00 0x00 0x00 0x00
0x17fbc0738: 0x00 0x00 0x00 0x00 0x00 0x00 0x00 0x00 0x00
```

CVL Flow Director Debug

How to debug FDIR rule ?

(E.g. pattern eth / ipv4 dst is 192.168.0.2 / end)

1. Check if Profile is created successfully and correctly
2. Check if dummy packet contains the correct training data
3. HW Write Back mechanism can confirm that packet is transmitted successfully.
4. Programming status report mechanism can confirm that packet is transmitted successfully.

FDIR set rule steps

1. Create Profile
2. Set rule through Flow Director TX/RX queue
 - a) Get free queue descriptor
 - b) Fill LAN+FDIR descriptor
 - c) Prepare FDIR dummy packet
 - d) Notify HW (Doorbell) and HW Write Back**
3. **Check programming status (Optional)**

CVL Flow Director Debug

How to debug FDIR rule ?

(E.g. pattern eth / ipv4 dst is 192.168.0.2 / end)

1. Check if Profile is created successfully and correctly
2. Check if dummy packet contains the correct training data
3. HW Write Back mechanism can confirm that packet is transmitted successfully.
4. Programming status report mechanism can confirm that packet is transmitted successfully.
5. Check Global/PF/VSI rule table space and size to make sure program succeed.

FDIR set rule steps

1. Create Profile
2. Set rule through Flow Director TX/RX queue
 - a) Get free queue descriptor
 - b) Fill LAN+FDIR descriptor
 - c) Prepare FDIR dummy packet
 - d) Notify HW (Doorbell) and HW Write Back
3. Check programming status (Optional)
4. **Check rule table space and size (Highly recommend)**

CVL Flow Director Debug

5. Check Global/PF/VSI rule table count increased to make sure program succeed.

At programming, the software can chose the space on which the filter is counted. If there is no space on the selected space, the programming request is rejected.

The software can track the number of programmed filters.

- Per PF: By *PFQF_FD_CNT.FD_GCNT* and *PFQF_FD_CNT.FD_BCNT* counters.
- Per VSI: By *VSIQF_FD_CNT.FD_GCNT* and *VSIQF_FD_CNT.FD_BCNT* counters.
- Globally: By the *GLQF_FD_CNT.FD_GCNT* and *GLQF_FD_CNT.FD_BCNT* counters.

Highly recommend:

- Important issue blocking DPDK PF FDIR POC and help a lot on IAVF FDIR POC.
- Almost necessary and sufficient condition of program success.
- Easy for debugging with little dependency on external tools.

CVL Flow Director Debug

How to debug FDIR rule ?

(E.g. pattern eth / ipv4 dst is 192.168.0.2 / end)

6. Send Test packets to LAN port and check if packets contains the desired input set value

E.g.

0xC0 0xA8 0x00 0x02 → 192.168.0.2

```
>>> hexdump(p_ipv4)
0000  00 00 00 00 01 03 A4 BF 01 51 27 CA 08 00 45 23 .....Q'...E#
0010  00 14 00 01 00 03 69 B9 C4 AB C0 A8 0B 0C C0 A8 .....i.....
0020  00 02 ..
>>> ls(p_ipv4)
dst      : DestMACField          = '00:00:00:00:01:03' (None)
src      : SourceMACField        = 'a4:bf:01:51:27:ca' (None)
type     : XShortEnumField       = 2048                (36864)
--
version  : BitField (4 bits)     = 4                   (4)
ihl      : BitField (4 bits)     = None                (None)
tos      : XByteField            = 35                  (0)
len      : ShortField            = None                (None)
id       : ShortField            = 1                   (1)
flags    : FlagsField (3 bits)   = <Flag 0 (>)         (<Flag 0 (>))
frag     : BitField (13 bits)    = 3                   (0)
ttl      : ByteField             = 105                 (64)
proto    : ByteEnumField         = 185                 (0)
chksum   : XShortField           = None                (None)
src      : SourceIPField         = '192.168.11.12'     (None)
dst      : DestIPField           = '192.168.0.2'       (None)
options  : PacketListField       = []                  ([])
```


CVL Flow Director Debug

How to debug FDIR rule ?

(E.g. pattern eth / ipv4 dst is 192.168.0.2 / end destined to queue 14)

6. Check if Test LAN packets contains the desired input set value

7. Check LAN rx packets descriptor and FDIR domain flags, For dpdk, you can see:

Received packet with **FDIR matched ID=XXX** is directed to queue 14

Notes:

Checking Queue ID is not the most reliable way since packets may be missed by FDIR and then distributed by RSS to the same queue index.

Please check FDIR flags like “FDIR matched”

```
port 0/queue 14: received 1 packets
src=A4:BF:01:51:27:CA - dst=00:00:00:00:01:03 - type=0x0800 - length=82 - nb_segs=1 - RSS hash=0x14ba313
6 - RSS queue=0xe - FDIR matched ID=0x0 - hw ptype: L2_ETHER L3_IPV4_EXT_UNKNOWN TUNNEL_GTPU INNER_L3_IPV4
_EXT_UNKNOWN INNER_L4_NONFRAG - sw ptype: L2_ETHER L3_IPV4 L4_UDP - l2_len=14 - l3_len=20 - l4_len=8 - V
XLAN packet: packet type =32913, Destination UDP port =2152, VNI = 1193046 - Receive queue=0xe
ol_flags: PKT_RX_RSS_HASH PKT_RX_FDIR PKT_RX_L4_CKSUM_GOOD PKT_RX_IP_CKSUM_GOOD PKT_RX_FDIR_ID PKT_RX_OUTER_L4_CKSUM_UNKNOWN
```

FDIR Intro

➤ Use case: PPPoE, GTP-U, etc.

flow create 0 ingress pattern eth / ipv4 / udp / gtpu teid is 0x12345678 /
gtp_psc qfi is 0x34 / end actions queue index 4 / end

Bits							
8	7	6	5	4	3	2	1
Version		PT	(*)	E	S	PN	
Message Type							
Length (1 st Octet)							
Length (2 nd Octet)							
Tunnel Endpoint Identifier (1 st Octet)							
Tunnel Endpoint Identifier (2 nd Octet)							
Tunnel Endpoint Identifier (3 rd Octet)							
Tunnel Endpoint Identifier (4 th Octet)							
Sequence Number (1 st Octet) ^{1) 4)}							
Sequence Number (2 nd Octet) ^{1) 4)}							
N-PDU Number ^{2) 4)}							
Next Extension Header Type ^{3) 4)}							

Bits								Num ber of Octet s
7	6	5	4	3	2	1	0	
PDU Type (=0)				Spare				1
PPP	RQI	QoS Flow Identifier						1
PPI				Spare				0 or 1
Padding								0-3

FDIR Use Case

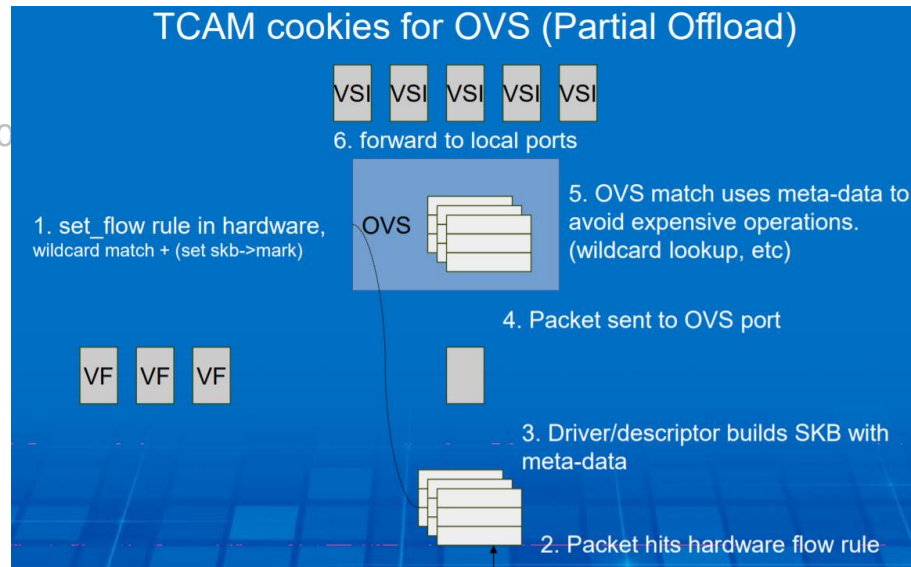
看一下。

- **Improve Memcached Performance:**
<https://www.intel.com/content/dam/www/public/us/en/documents/white-papers/intel-ethernet-flow-director.pdf>
- **Improve Redis Performance:**
<https://www.intel.com/content/www/us/en/architecture-and-technology/ethernet/application-device-queues-with-redis-brief.html>


FDIR Use Case

➤ Use case: OVS Partial Offload:

1. Set_flow rule in HW
flow create 0 ingress pattern eth / ipv4 / udp dst is 80 / end actions mark id 0x1234 end
2. Packet hits hw flow rule and mark **desc**
with **FD_ID** field set to 0x1234
3. Driver builds SKB with marks in step 2
4. Packet sent to OVS Port
5. **OVS match uses mark id 0x1234 to avoid expensive lookup**
6. Forward to local ports



Agenda



AVF Introduction

RSS Introduction

FDIR Introduction

Virtual Channel Introduction

Virtual Channel Intro

➤ What is Virtual Channel:

VF PF virtual communication channel

avf channel 分为base mode和advance mode。

➤ Why Virtual Channel?

- Resource negotiation: RSS, VLAN, MAC, QUEUE, IRQ...
- HW management: RESET, EVENT...

Virtual Channel Intro

➤ Supported Virtual Channel Command:

```
VIRTCHNL_OP_UNKNOWN = 0,  
VIRTCHNL_OP_VERSION = 1, /* must be 1 */  
VIRTCHNL_OP_RESET_VF = 2,  
VIRTCHNL_OP_GET_VF_RESOURCES = 3,  
VIRTCHNL_OP_CONFIG_TX_QUEUE = 4,  
VIRTCHNL_OP_CONFIG_RX_QUEUE = 5,  
VIRTCHNL_OP_CONFIG_VSI_QUEUES = 6,  
VIRTCHNL_OP_CONFIG_IRQ_MAP = 7,  
VIRTCHNL_OP_ENABLE_QUEUES = 8,  
VIRTCHNL_OP_DISABLE_QUEUES = 9,  
VIRTCHNL_OP_ADD_ETH_ADDR = 10,  
VIRTCHNL_OP_DEL_ETH_ADDR = 11,  
VIRTCHNL_OP_ADD_VLAN = 12,  
VIRTCHNL_OP_DEL_VLAN = 13,  
VIRTCHNL_OP_CONFIG_PROMISCUOUS_MODE = 14,  
VIRTCHNL_OP_GET_STATS = 15,  
VIRTCHNL_OP_RSVD = 16,  
VIRTCHNL_OP_EVENT = 17, /* must be 17 */
```

base mode

```
VIRTCHNL_OP_CONFIG_RSS_KEY = 23,  
VIRTCHNL_OP_CONFIG_RSS_LUT = 24,  
VIRTCHNL_OP_GET_RSS_HENA_CAPS = 25,  
VIRTCHNL_OP_SET_RSS_HENA = 26,  
VIRTCHNL_OP_ENABLE_VLAN_STRIPPING = 27,  
VIRTCHNL_OP_DISABLE_VLAN_STRIPPING = 28,  
VIRTCHNL_OP_REQUEST_QUEUES = 29,  
VIRTCHNL_OP_ENABLE_CHANNELS = 30,  
VIRTCHNL_OP_DISABLE_CHANNELS = 31,  
VIRTCHNL_OP_ADD_CLOUD_FILTER = 32,  
VIRTCHNL_OP_DEL_CLOUD_FILTER = 33,  
VIRTCHNL_IPSEC  
VIRTCHNL_OP_INLINE_IPSEC = 34,  
  
/* opcode 34 is reserved */  
/* VIRTCHNL_IPSEC */  
DCF_SUPPORT  
VIRTCHNL_OP_DCF_CMD_DESC = 39,  
VIRTCHNL_OP_DCF_CMD_BUFF = 40,
```

```
DCF_SUPPORT  
VIRTCHNL_OP_DCF_CMD_DESC = 39,  
VIRTCHNL_OP_DCF_CMD_BUFF = 40,  
VIRTCHNL_OP_DCF_DISABLE = 41,  
VIRTCHNL_OP_DCF_GET_VSI_MAP = 42,  
VIRTCHNL_OP_DCF_GET_PKG_INFO = 43,  
  
/* opcodes 39, 40, 41, 42 and 43 are reserved */  
/* DCF_SUPPORT */  
ADV_AVF_SUPPORT  
VIRTCHNL_OP_GET_SUPPORTED_RXDIRS = 44,  
VIRTCHNL_OP_ADD_RSS_CFG = 45,  
VIRTCHNL_OP_DEL_RSS_CFG = 46,  
VIRTCHNL_OP_ADD_FDIR_FILTER = 47,  
VIRTCHNL_OP_DEL_FDIR_FILTER = 48,  
VIRTCHNL_OP_QUERY_FDIR_FILTER = 49,  
  
/* opcode 44, 45, 46, 47, 48 and 49 are reserved */  
/* ADV_AVF_SUPPORT */
```

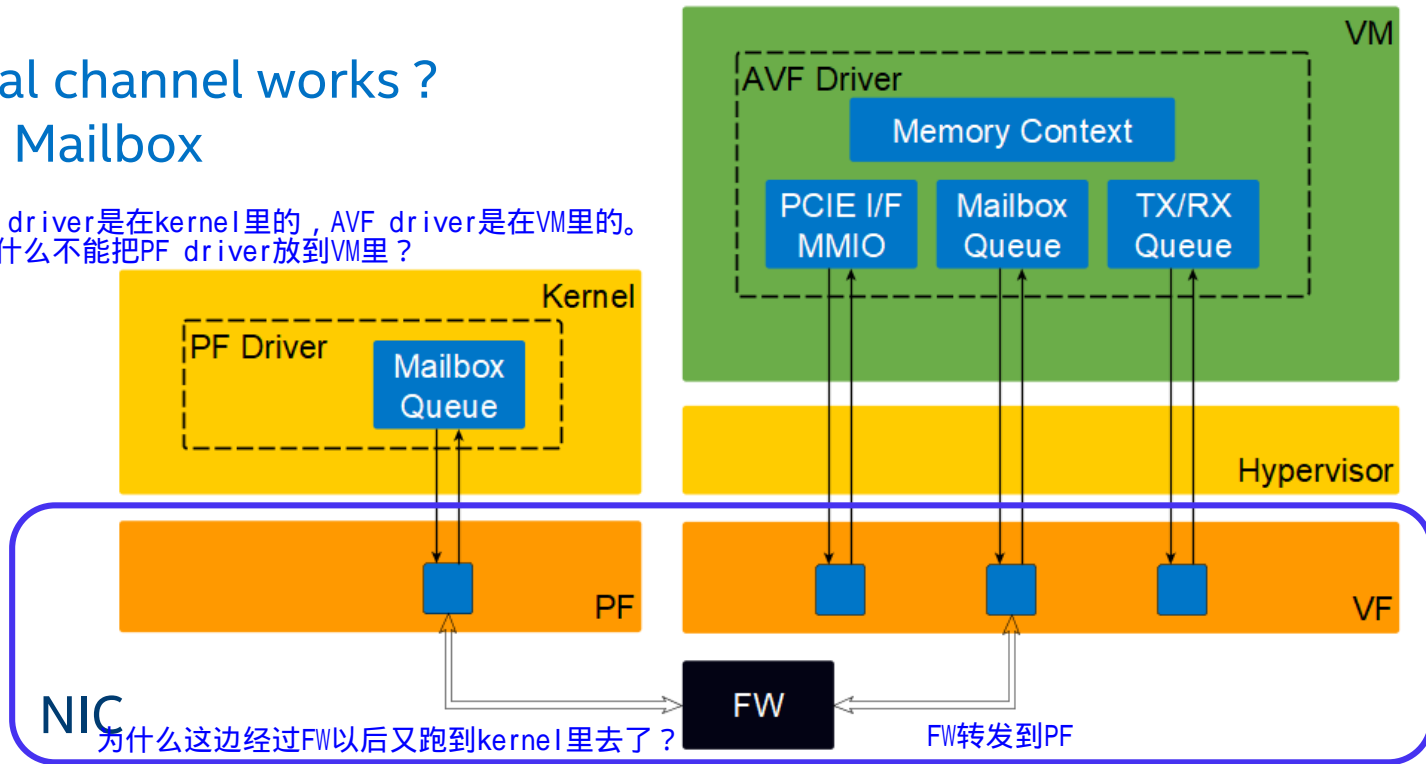
DCF, AVF相关 (添加RSS config, 添加flow director filter)

Virtual Channel Intro

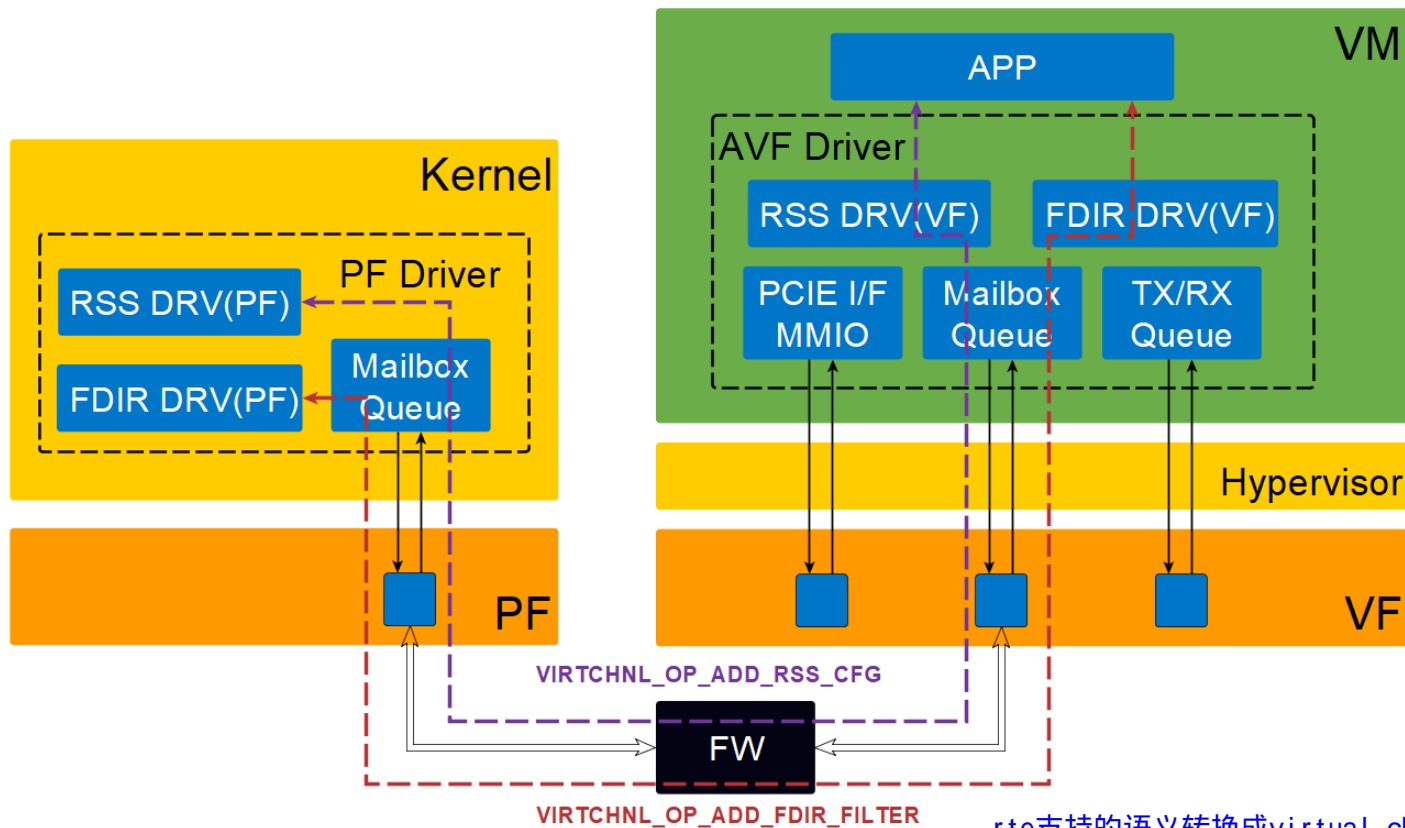
和PF通信的三种方式：PCIE MMIO, Mailbox, TX/RX

➤ How virtual channel works ? Hardware Mailbox

PF driver是在kernel里的，AVF driver是在VM里的。
为什么不能把PF driver放到VM里？



Virtchnl for RSS/FDIR design



hackson: virtual channel的workload

rte支持的语义转换成virtual channel支持的格式。

Q & A